

Learning to Understand: Identifying Interactions via the Möbius Transform

Berkeley
UNIVERSITY OF CALIFORNIA

Justin S. Kang¹ Yigit E. Erginbas¹ Landon Butler¹ Ramtin Pedarsani² Kannan Ramchandran¹

UCSB

¹UC Berkeley ²UC Santa Barbara

Problem

Deep learning models are getting better, but not any easier to understand.

- A popular approach for building explanations of models involves looking at first-order approximations, like the well known Shapley Value.
- First order models can miss important structures critical for explanation.
- Example: A sentiment analysis LLM trained on the IMDB dataset:




Figure 1: Presented are 1st, 2nd and 3rd order Möbius coefficients. While never and fails have negative sentiments, combined they are strongly positive. In the second row, the word never is deleted, changing overall sentiment. The Shapley values SV(·) are less informative.

- The word “never” has a negative first-order sentiment, but is involved in critical second order interactions, making its net effect positive.

The Möbius Transform

- The model for higher order interactions is called the Möbius Transform:

$$\text{Inverse: } f(\mathbf{m}) = \sum_{\mathbf{k} \leq \mathbf{m}} F(\mathbf{k}), \quad \text{Forward: } F(\mathbf{k}) = \sum_{\mathbf{m} \leq \mathbf{k}} (-1)^{\mathbf{k}^T (\mathbf{k} - \mathbf{m})} f(\mathbf{m})$$

Naïve computation is exponential in number of features n .

- Compare with the Shapley Values SV(·) and Banzhaf Values BZ(·):

$$SV(i) = \sum_{\mathbf{k}: k_i=1} \frac{1}{|\mathbf{k}|} F(\mathbf{k}), \quad BZ(i) = \sum_{\mathbf{k}: k_i=1} \frac{1}{2^{|\mathbf{k}|-1}} F(\mathbf{k}).$$

- A small number of interactions dominate the function overall.




Figure 2: $F(\mathbf{k})$ generally has a sparse structure. The functions are well-approximated with only a small number of coefficients (sparsity), and these coefficients also have small $|\mathbf{k}|$ (low degree). Can we compute the Möbius transform more efficiently under these settings?


The Algorithm

Step 1: Aliasing Informed Masking Design

- Construct the function u from samples of f with $b \ll n$, and take the Transform of u , denoted U in $b2^b$ time:

$$u_c(\ell) = f(\mathbf{H}_c^T \ell) \quad \forall \ell \in \mathbb{Z}_2^b \iff U_c(\mathbf{j}) = \sum_{\mathbf{H}, \mathbf{k}=\mathbf{j}} F(\mathbf{k}) \quad \forall \mathbf{j} \in \mathbb{Z}_2^b.$$

- Aliasing effectively hashes the coefficients $F(\mathbf{k})$ into one of 2^b bins:



- The singleton coefficients can be detected, and their \mathbf{k} index identified.

Step 2: Identifying Interactions via Group Testing

- The key to identifying a singletons is to construct “delayed” versions of u :

$$u_{cp}(\ell) = f(\mathbf{H}_c^T \ell + \mathbf{d}_p) \iff U_c(\mathbf{j}) = \sum_{\substack{\mathbf{H}, \mathbf{k}=\mathbf{j} \\ \mathbf{k} \leq \mathbf{d}_p}} F(\mathbf{k}).$$

- A “delay” is a membership test on \mathbf{k} . Repeating, we construct $\mathbf{y} = \mathbf{D}\mathbf{k}$.
- When \mathbf{k} is arbitrary we take $\mathbf{D} = \mathbf{I}$, and require n delays \mathbf{d}_p .
- When $|\mathbf{k}| < t$ for some t , we choose \mathbf{D} as a group testing matrix:


$\mathbf{k}_1 =$	Her	acting	never	fails	to	impress	\mathbf{y}
	0	0	0	1	1	1	0
$\mathbf{D} =$	0	1	(1)	0	0	1	1
	1	0	(1)	0	1	0	1

Decode $\mathbf{y} \longrightarrow \mathbf{k}_1$

- Theory says we only require $O(t \log(n))$ delays to ensure recovery.


Step 3: Message Passing to Resolve Collisions

- Defines a bipartite graph connecting the non-zero $F(\mathbf{k})$ and U .
- Use a message passing algorithm (peeling decoder) to resolve multitons. This is inspired by sparse graph codes for robust communication.



- Choosing \mathbf{H}, \mathbf{D} correctly ensures we are likely to peel all non-zero $F(\mathbf{k})$.
- Density evolution theory can prove the performance of the algorithm.

Overview



We design masking patterns according to a group testing design, and perform inference of the masked inputs. If needed, the output is converted to a scalar, and the output is used to compute the Möbius Transform.

Our algorithm is non-adaptive and has rigorous performance guarantees.

Theorems

- (Sparse) With K non-zero interactions among all 2^n interaction, our algorithm exactly computes the Möbius transform $F(\mathbf{k})$ in $O(Kn)$ samples and $O(Kn^2)$ time with probability $1 - O(1/K)$.
- (Sparse, Low Degree) When there are K non-zero interactions all with $|\mathbf{k}| \leq t$, our algorithm computes the Möbius transform in $O(Kt \log(n))$ samples and $O(K \text{poly}(n))$ time with probability $1 - O(1/K)$, even under the presence of noise at any fixed SNR.

Experiments




Figure: (a) Sample complexity of our algorithm. Clear phase transition, with the threshold scaling linearly in n is visible. (b) Shows our algorithm under a noise model where $U(\mathbf{j})$ are corrupted by Gaussian noise at different SNR.




Figure: Using only a small number of coefficients (sparsity), the Möbius transform computed by our method outperforms first order methods in faithfulness (R^2) to the underlying network. The gap is larger in problems with non-linear feature relationships.

Further Reading

- Kang JS, et al. "Learning to Understand: Identifying Interactions via the Möbius Transform". NeurIPS (2024).
- Erginbas YE, Kang JS et al.. "Efficiently Computing Sparse Fourier Transforms of q -ary Functions." IEEE ISIT (2023).

